

CSI 4107

Assignment 2

Neural information retrieval system

Nick Ivanov, 300137667
Bradley Liu, 300125419
Kevin Le, 300053306

University of Ottawa

Link to github repository: <https://github.com/nivanov8/csi4107-assignment2>

1. How to run:

- a. Run: ``pip install -r requirements.txt``
 - b. In the root directory run:
 - i. ``python main.py``
- Note: If you run into an error, make sure Windows Long Path is enabled. You might have to check in your Windows Registry Editor to enable it in order to download these packages.

2. Functionality:

- a. Preprocessing:

The tokenization for this assignment was done the exact same as Assignment 1, however, we will explain it again in the next paragraphs as it is a big role in information retrieval.

Tokenization: The system tokenizes documents into individual words or tokens, removing punctuation and special characters.

Stop Word Removal: Common stop words such as "and", "the", and "is" are removed to reduce noise in the text.

The preprocessing functionality provided involves several steps to prepare textual data for indexing and retrieval. First, tokenization is performed using the `tokenize` method, which splits text into individual words or tokens. To ensure consistency, the text is converted to lowercase, and regular expressions are employed to remove punctuation, special characters, and underscores. Additionally, apostrophes are normalized to handle cases like double apostrophes.

The `preprocess()` method iterates through each document file in the specified folder. Documents are split into individual segments using the `<DOC>` tag as a delimiter. For each segment, the `<DOCNO>` tag is extracted to identify the document number, while the `<TEXT>` tag is used to obtain the textual content. This content is then tokenized using the `tokenize()` method, and the resulting tokens are stored in a list.

The tokenized text from each document segment is joined back into a single string and added to the corpus, representing the entire collection of documents. The method returns three items: a list

of tokens, the corpus of documents, and the mapping of text to document numbers. The return values for the preprocessing used for tf-idf is a bit different due to used a different library.

Note that we also preprocess the documents held in the collection folder differently from the text-queries.txt. As their delimiters are different, in similar fashion we strip the whitespace and build the queries by line and return the list of queries.

b. BM25 retrieval

As mentioned per the project outline, BM25 from assignment 1 was used to retrieve the first 1000 documents to eliminate the need to use a neural model on a large corpus of documents. We used the same BM25 model from assignment 1 (BM25OKAPI) to retrieve the initial subset of documents. This retrieval takes place in the bm25 folder of the project.

c. Neural reranking

After the initial retrieval of the 1000 documents for each query using BM25, we then used 2 models to rerank the results neurally. The first model used was GIST large embedding v0 (model card: <https://huggingface.co/avsolatorio/GIST-large-Embedding-v0>). This model was selected due to its good performance on the MTEB (Massive Text Embedding Benchmark) Leaderboard on Huggingface. This model is primarily used for sentence similarity and we thought this would be appropriate to use as the nature of embeddings and similarities of embeddings is to capture sentence similarities. This model did the best out of the 2 used. The second model used was MxBai large embedding v1 (model card: <https://huggingface.co/mixedbread-ai/mxbai-embed-large-v1>). This was an instruction based model so we had to reformulate the queries such that we can better assist the model in generating embeddings as this was how the model was trained. This model was a close alternative for reranking.

d. Evaluation:

Mean Average Precision (MAP): We compute the MAP score using trec_eval, a standard tool for evaluating information retrieval systems. MAP measures the average precision of the system across multiple queries.

Precision and Recall: We also calculate precision and recall metrics to assess the effectiveness of the retrieval system in retrieving relevant documents.

3. Results

The results obtained from this assignment concluded that using a neural model to rerank the initial documents obtained only slightly increased the resulting MAP score. In Assignment 1 (using purely BM25) we managed to achieve a MAP score of 0.3411. Using our best neural model, GIST Large Embedding V0, we obtained a MAP score of 0.3450, a 1.1% increase in performance. Although not provided, the MAP score obtained for using the MxBai model was 0.3120 and as previously mentioned this was a close alternative to the better performing GIST model.

4. Algorithms and Optimizations

a. Preprocessing

We used regex matching to separate each document and extract information like the title and the text. We then ran the documents through a tokenizer that would transform the text into all lower case, remove stop words, stem and remove any special characters. We would then return a list of tokens to be used for indexing and other information such as the corpus itself and all the doc numbers for easier computations later down the line. This was the exact same technique we used in assignment 1

b. Retrieval and Ranking

We only retrieved a subset of documents from the whole corpus using bm25 first. This drastically reduces the computation time as calculating vectors for the text can take a long time even for 1000 documents. This method helped reduce the computational power and sorting the documents according to similarities scores was also faster due to only 1000 documents being present.

5. 10 Results from query 3 and 20

a. Query 3

3 Q0 AP880419-0133 1 0.6601 BM25
3 Q0 AP880920-0017 2 0.619 BM25
3 Q0 AP880518-0053 3 0.6029 BM25
3 Q0 AP880523-0108 4 0.5272 BM25
3 Q0 AP880609-0027 5 0.5228 BM25

3 Q0 AP880628-0074 6 0.5206 BM25
3 Q0 AP880609-0025 7 0.5189 BM25
3 Q0 AP880701-0160 8 0.5186 BM25
3 Q0 AP880511-0043 9 0.5168 BM25
3 Q0 AP880829-0244 10 0.5087 BM25

b. Query 20

20 Q0 AP881111-0092 1 0.7423 BM25
20 Q0 AP881110-0035 2 0.7239 BM25
20 Q0 AP881122-0171 3 0.7235 BM25
20 Q0 AP881123-0037 4 0.6942 BM25
20 Q0 AP880627-0239 5 0.6887 BM25
20 Q0 AP880323-0094 6 0.5632 BM25
20 Q0 AP880314-0099 7 0.5528 BM25
20 Q0 AP881010-0289 8 0.5478 BM25
20 Q0 AP881103-0013 9 0.5416 BM25
20 Q0 AP880314-0323 10 0.537 BM25

Note: These results are straight from the Results.txt provided. BM25 and Q0 are just placeholder names and do not matter.

6. MAP and P@10

a. MAP Score

```
nickivanov@Nicks-MacBook-Pro-2 trec_eval-9.0.7 % ./trec_eval -q -m map qrels1-50ap.txt Results.txt
map      1      0.2125
map     10      0.4280
map     11      0.7846
map     12      0.4287
map     13      0.7562
map     14      0.2605
map     15      0.0898
map     16      0.6022
map     17      0.3906
map     18      0.4097
map     19      0.2773
map      2      0.5215
map     20      0.9167
map     21      0.0974
map     22      0.1994
map     23      0.6931
map     24      0.4870
map     25      0.2693
map     26      0.0635
map     27      0.3703
map     28      0.4723
map     29      0.3791
map      3      0.3881
map     30      0.1855
map     31      0.3068
map     32      0.3113
map     33      0.4050
map     34      0.0481
map     35      0.5084
map     36      0.1626
map     37      0.3185
map     38      0.1752
map     39      0.3123
map      4      0.5665
map     40      0.3079
map     41      0.3175
map     42      0.1651
map     43      0.5868
map     44      0.1093
map     45      0.3196
map     46      0.3034
map     47      0.2697
map     48      0.3717
map     49      0.1000
map      5      0.0449
map     50      0.4577
map      6      0.5944
map      7      0.2514
map      8      0.0034
map      9      0.2500
map     all     0.3450
```

As seen from the screenshot at the very last line, the MAP score is 0.3450

b. $P@10$

```

nickivanov@Nicks-MacBook-Pro-2 trec_eval-9.0.7 % ./trec_eval qrels1-50ap.txt Results.txt
runid          all      BM25
num_q          all      50
num_ret        all      49533
num_rel        all      2099
num_rel_ret    all      1712
map            all      0.3450
gm_map         all      0.2651
Rprec          all      0.3379
bpref          all      0.3826
recip_rank     all      0.7512
iprec_at_recall_0.00 all    0.7846
iprec_at_recall_0.10 all    0.6506
iprec_at_recall_0.20 all    0.5530
iprec_at_recall_0.30 all    0.4650
iprec_at_recall_0.40 all    0.4098
iprec_at_recall_0.50 all    0.3530
iprec_at_recall_0.60 all    0.2793
iprec_at_recall_0.70 all    0.2157
iprec_at_recall_0.80 all    0.1554
iprec_at_recall_0.90 all    0.0845
iprec_at_recall_1.00 all    0.0418
P_5            all      0.5120
P_10           all      0.4500
P_15           all      0.4147
P_20           all      0.3900
P_30           all      0.3407
P_100          all      0.2048
P_200          all      0.1276
P_500          all      0.0623
P_1000         all      0.0342

```

As seen from the screenshot, the corresponding P@10 is 0.4500

Note: these scores are based of the GIST Large Embeddings V0 model as this model gave the best results.

7. Division of Tasks:

Kevin: Preprocessing, Evaluation, Implementation of MxBai model, Report

Nick: Preprocessing, Retrieval, Evaluation, Implementation of GIST model, Report

Bradley: Evaluation, Retrieval, Experimental evaluation of other models, Report